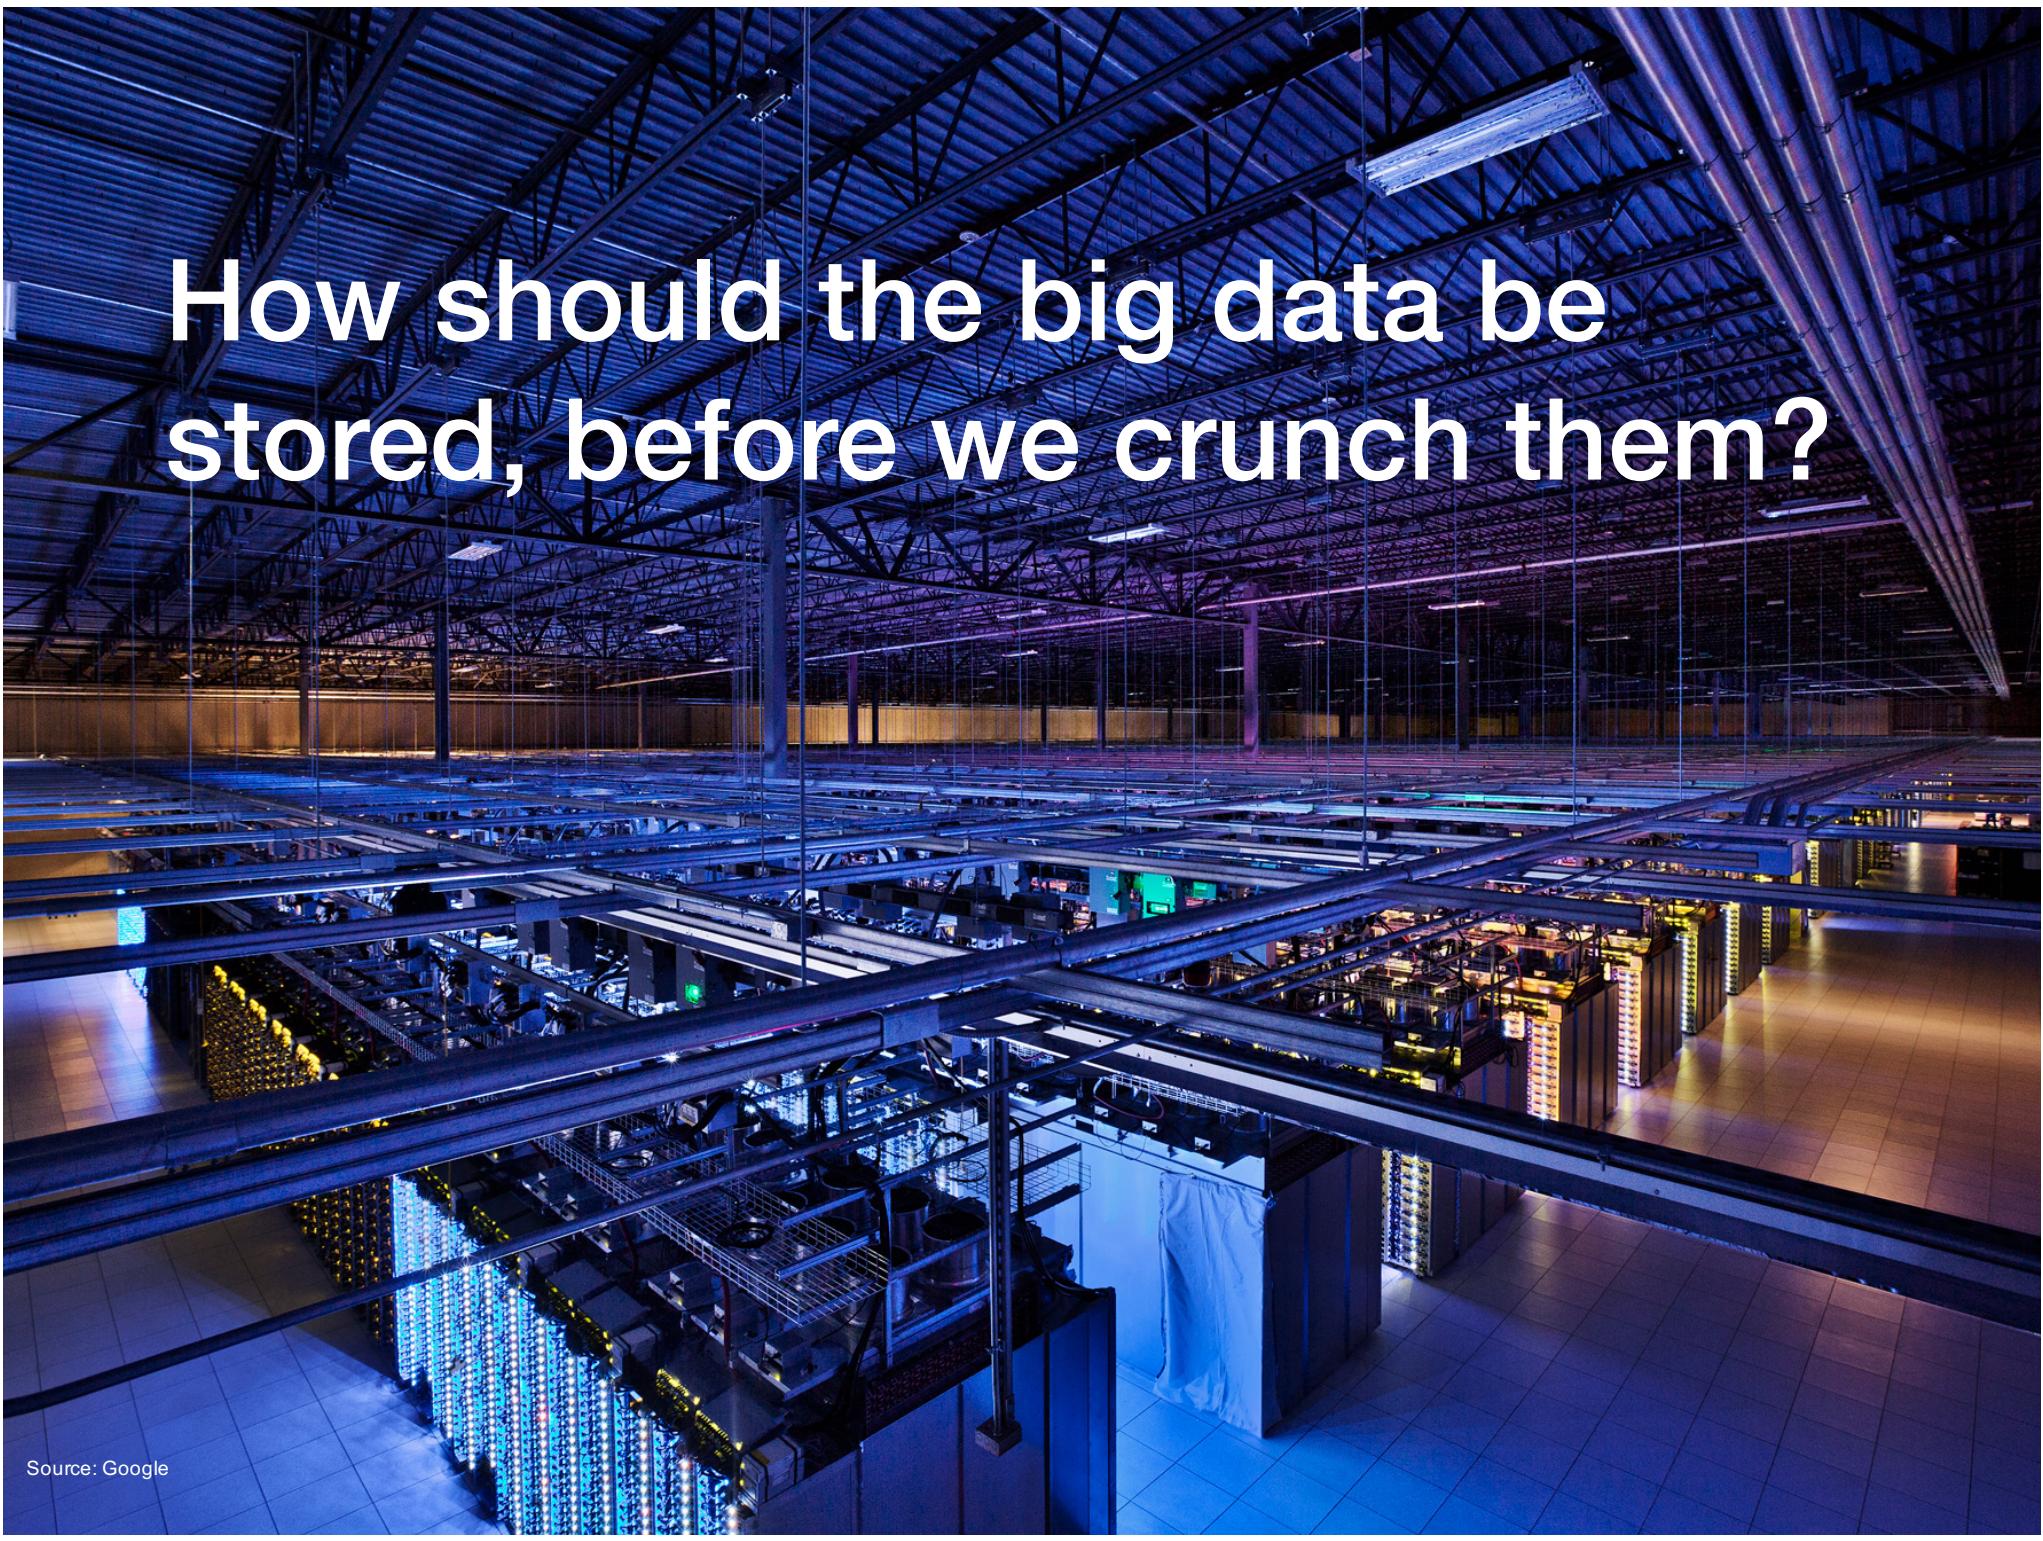


# Advanced Cloud Computing Distributed Storage System

---

Wei Wang  
CSE@HKUST  
Spring 2022





How should the big data be stored, before we crunch them?

Can we just have a BIG disk?

# Motivating app: Web search

---

Crawl the whole web

Store it all on “one **BIG** disk”

Process users’ searches on “one **BIG** CPU”

**Does it scale?**



# I/O is a bottleneck

---

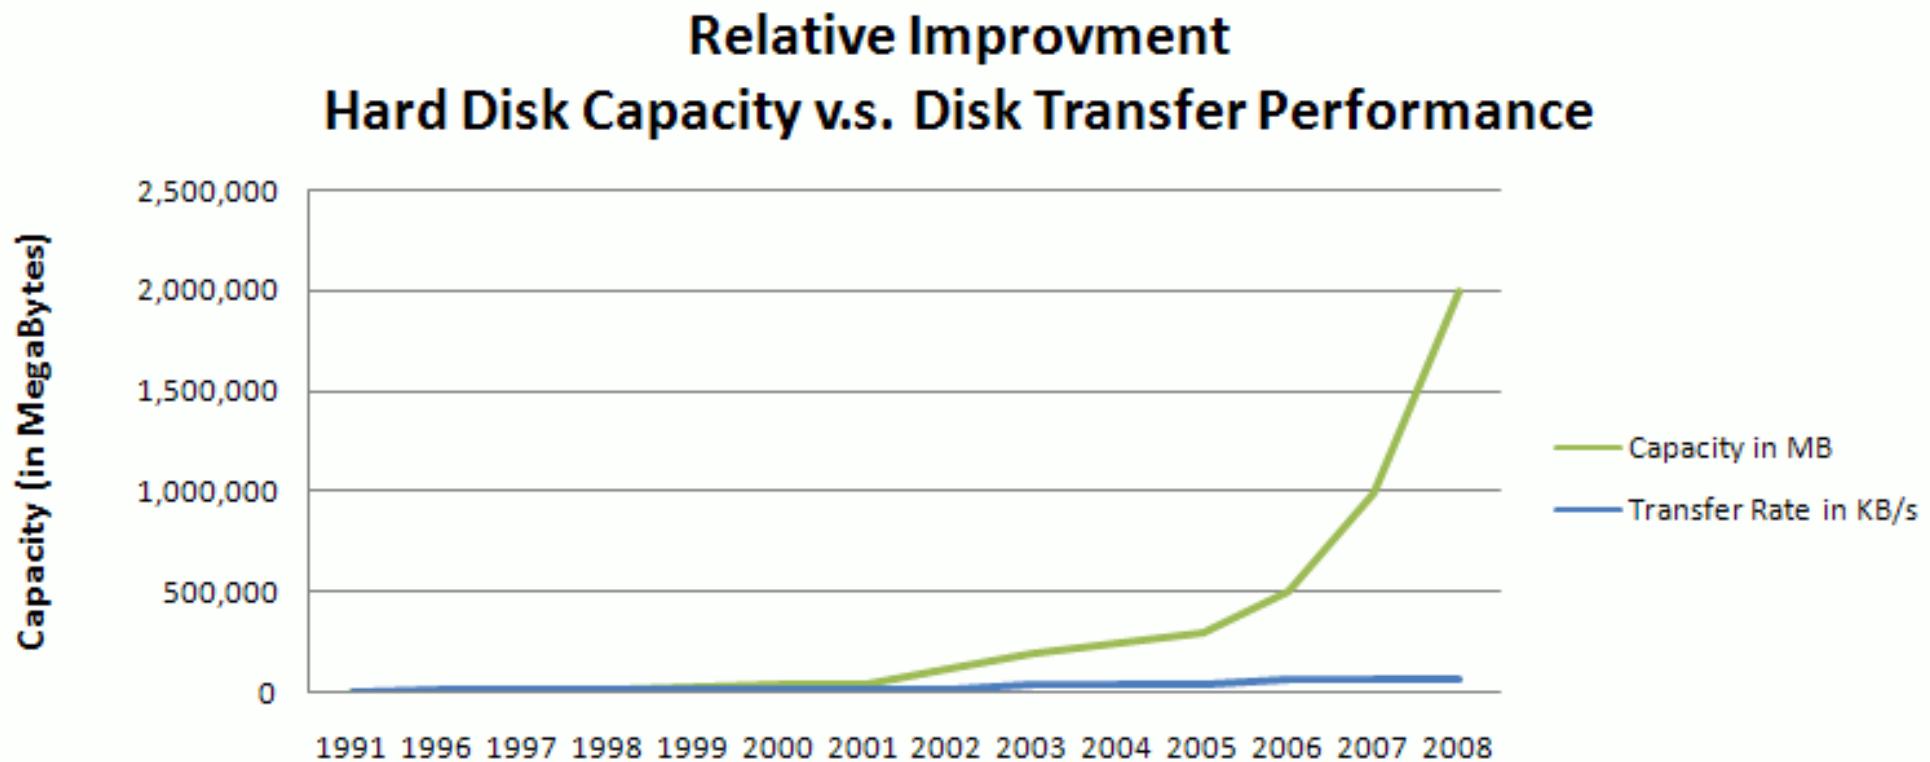
Suppose we have crawled 100 Tb worth of data

State-of-the-art 7,200 rpm SATA drive has 3 Gbps I/O

It takes  $100 \text{ Tb} / 3 \text{ Gbps} = \mathbf{9.3 \text{ hours}}$  to scan through the entire data!

# I/O lags far behind

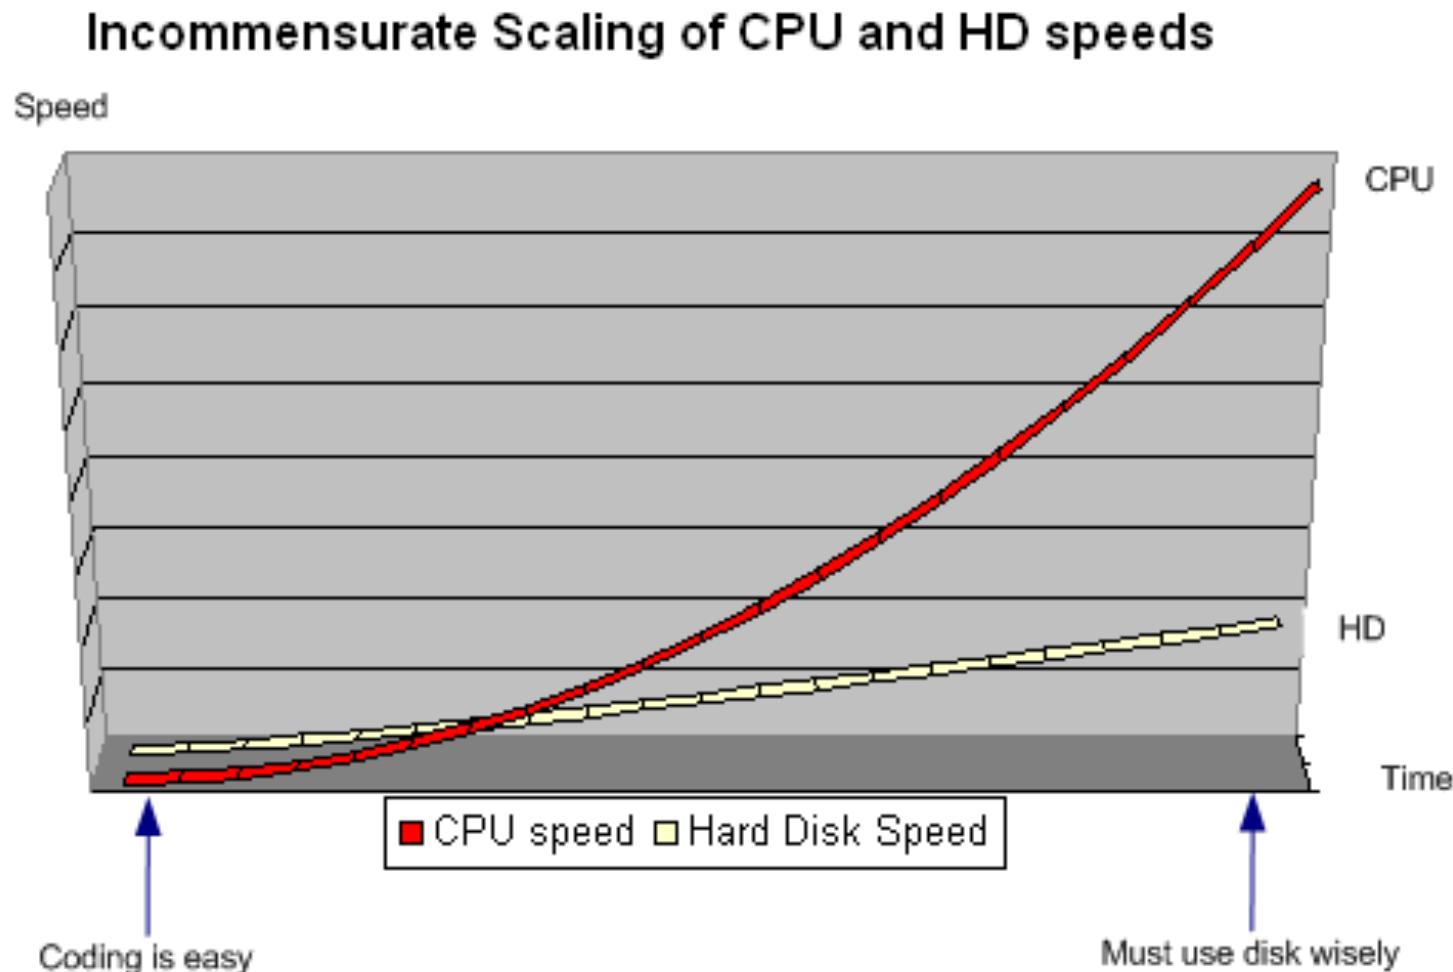
---



Source: R1Soft, <http://wiki.r1soft.com/pages/viewpage.action?pageId=3016608>

# I/O lags far behind

---



Source: R1Soft, <http://wiki.r1soft.com/pages/viewpage.action?pageId=3016608>

# Other issues

---

Building a high-end supercomputer is much, much costly

Storing all data in one place adds the risk of hardware failures

- ▶ putting all eggs in one basket!

Is there a way out?

Google's answer: scale “out”,  
not “up”!

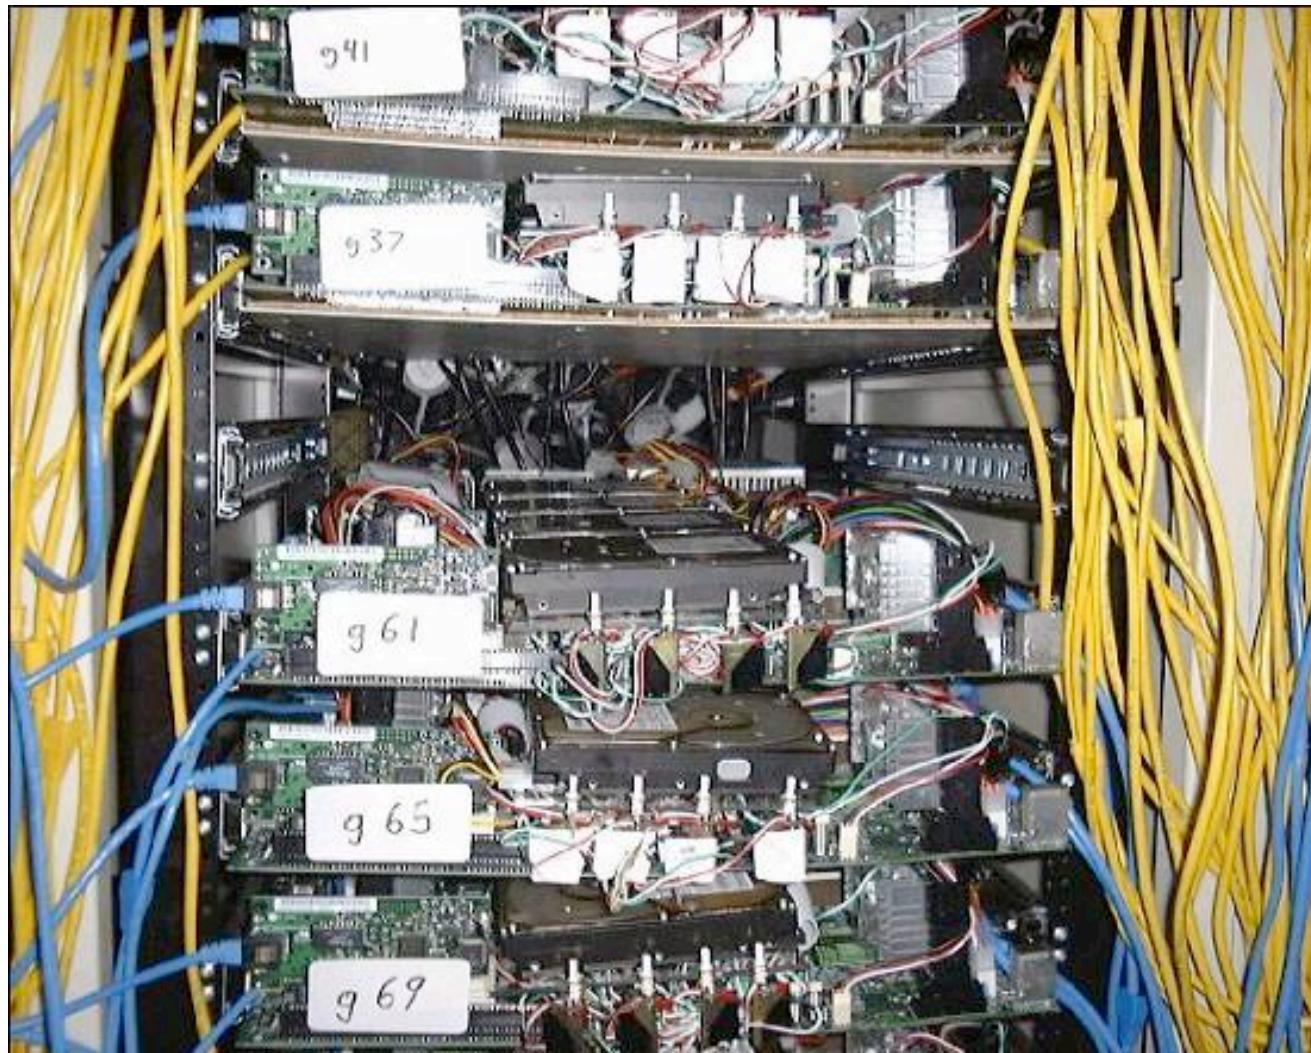
# “Google” circa 1997 (in Stanford)

---



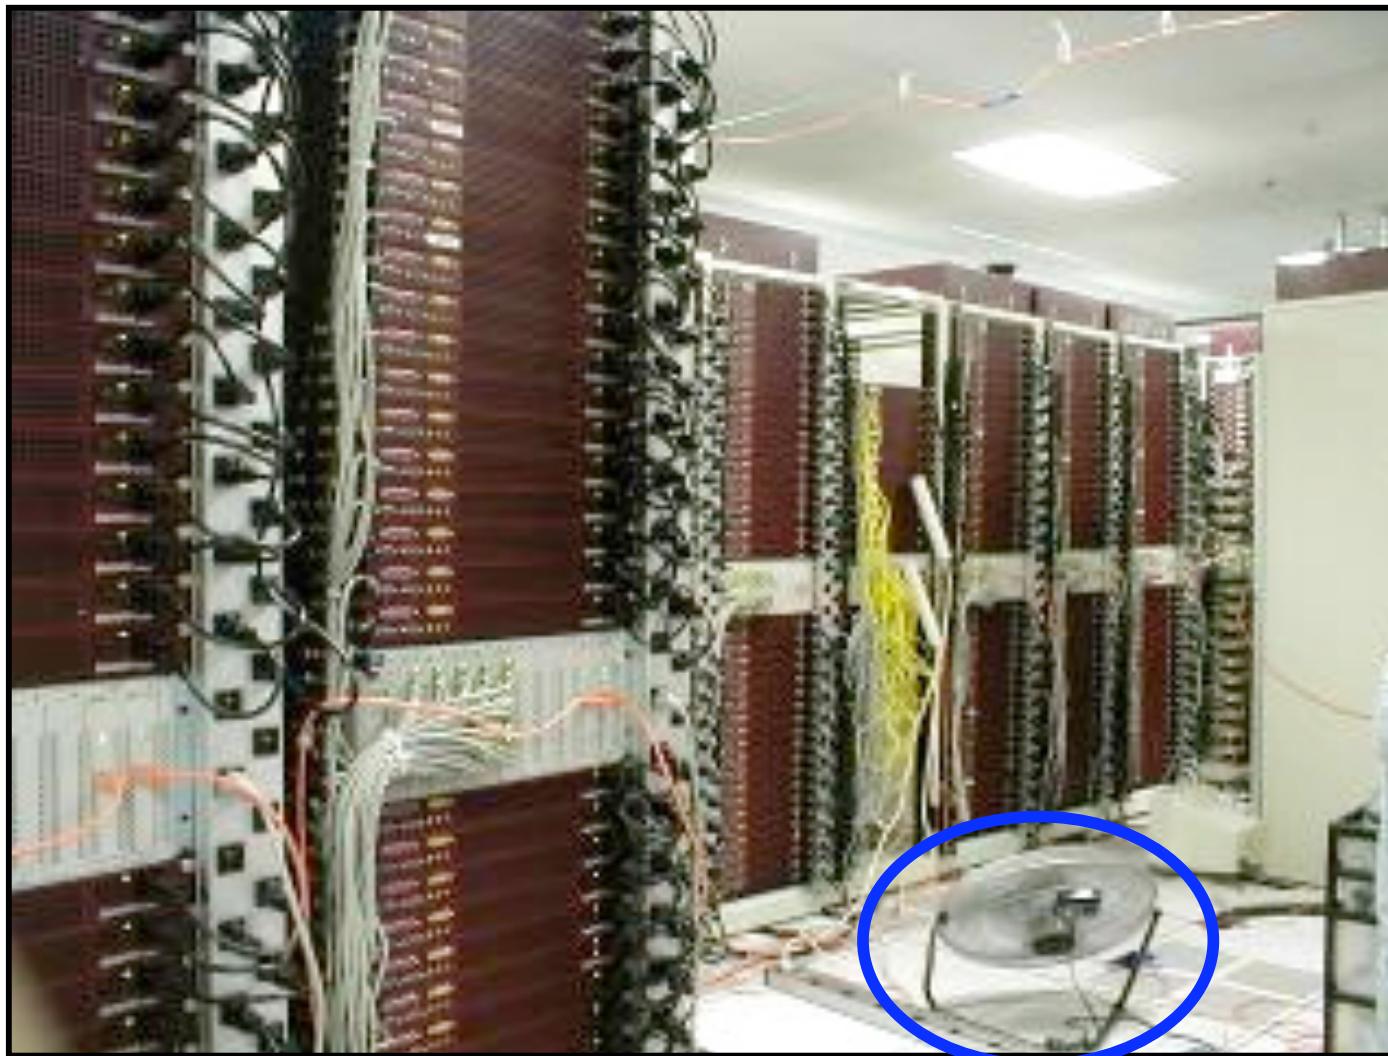
# Google (circa 1999)

---



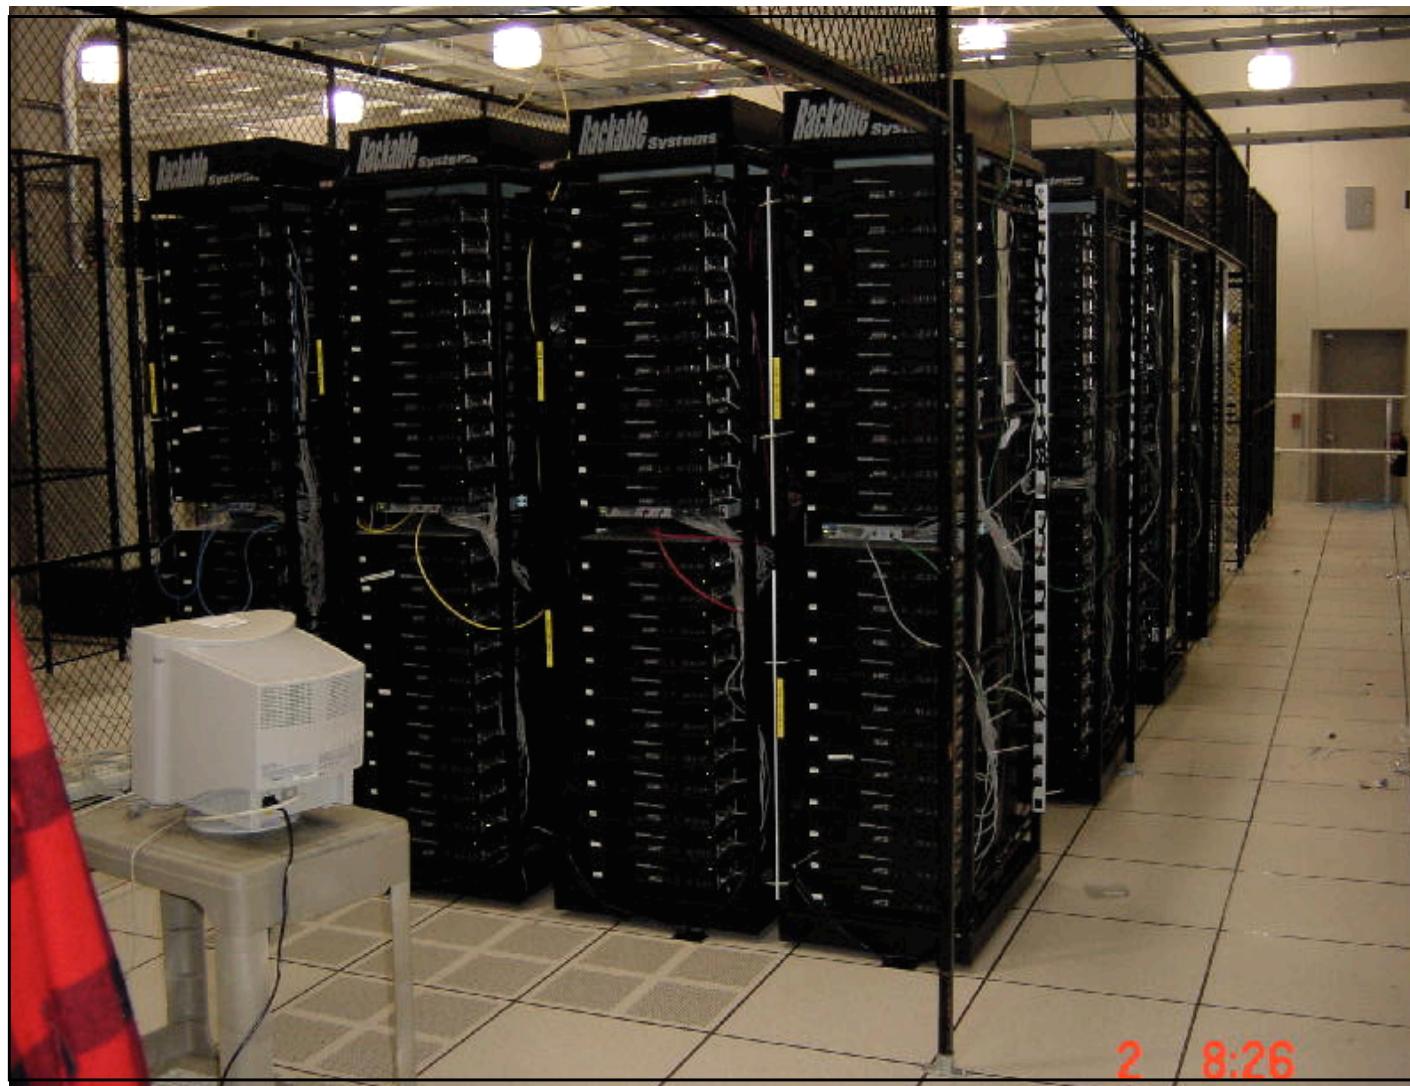
# Google datacenter (circa 2000)

---



# Google datacenter (circa 2001)

---



# Scale “out”, not “up”!

---

Lots of cheap, commodity PCs, each with disk and CPU

- ▶ 100s to 1000s of PCs in cluster in early days

High aggregate storage capacity

- ▶ No costly “big disk”

Spread search processing across many machines

- ▶ High I/O bandwidth, proportional to the # of machines
- ▶ Parallelize data processing

# Scale “out”, not “up”!

---

Suppose we have crawled 100 Tb worth of data and stored in a 1000-node cluster

- ▶ State-of-the-art 7,200 rpm SATA drive has 3 Gbps I/O
- ▶ 1000 nodes provide 3 Tbps aggregated I/O, 1000x faster than the BIG disk

It takes  $100 / 3 = \text{33 seconds}$  to scan through the entire data!

A cool idea! But wait...

# The joys of real HW in the 1st year

---

~0.5 **overheating** (power down most machines in <5 mins, ~1-2 days to recover)

~1 **PDU failure** (~500-1000 machines powered down, ~6 hours)

~1 **rack-move** (~500-1000 machines powered down, ~6 hours)

~1 **network rewiring** (rolling ~5% of machines down over 2-day span)

~20 **rack failures** (40-80 machines instantly disappear, 1-6 hours to get back)

~5 **racks go wonky** (40-80 machines see 50% packet loss)

~8 **network maintenance** (4 might cause ~30-min connectivity loss)

~1000 **individual machine failures**, ~thousands of **hard drive failures**, **slow disks**, **bad memory**, **misconfigured machines**, **flaky machines**, etc.

The list goes on...

# Implications

---

Stuff breaks

- ▶ If you have one server, it may stay up 3 years (1,000 days)
- ▶ If you have 10k servers, expect to lose 10 a day

“Ultra-reliable” hardware doesn’t really help

- ▶ At large scales, super-fancy reliable hardwares still fails, albeit less frequently
  - ▶ software still needs to be fault-tolerant
  - ▶ commodity machines w/o fancy h/w give better perf/\$

**Reliability has to come from  
the software!**

# GFS: The Google File System

A highly reliable storage system built  
atop highly unreliable hardwares

# Outline

---

Target environment

Design decisions

General architecture

File read and write

Fault tolerance

Measurements

Conclusions

# Target environment

---

Thousands of computers

Distributed

- ▶ Computers have their own disks, and the file system spans those disks

Failures are the **norm**

- ▶ Disks, networks, processors, power supplies, application software, OS software, human errors

# Target environment

---

Files are **huge**, but **not many**

- ▶ >100M, usually multi-gigabyte
- ▶ a few million files

**Write-once, read-many**

- ▶ Files are mutated by **appending**
- ▶ Once written, files are typically **only read**
- ▶ Large streaming reads and small random reads are typical

# Target environment

---

I/O bandwidth is more important than latency

- ▶ Suitable for **batch** processing and log analytics

It's helpful if the file system provides synchronization for  
**concurrent appends**

# GFS Design Decisions

# File stores as/divided into chunks

---

Large chunk size: 64 MB

- ▶ a file smaller than a single chunk does not occupy a full chunk's worth of storage

Why large chunks?

- ▶ minimizes the cost of disk seeks
- ▶ pushes up file transfer rate to the disk transfer rate

$$T_{\text{File transfer}} = T_{\text{disk transfer}} + T_{\text{disk seek}}$$

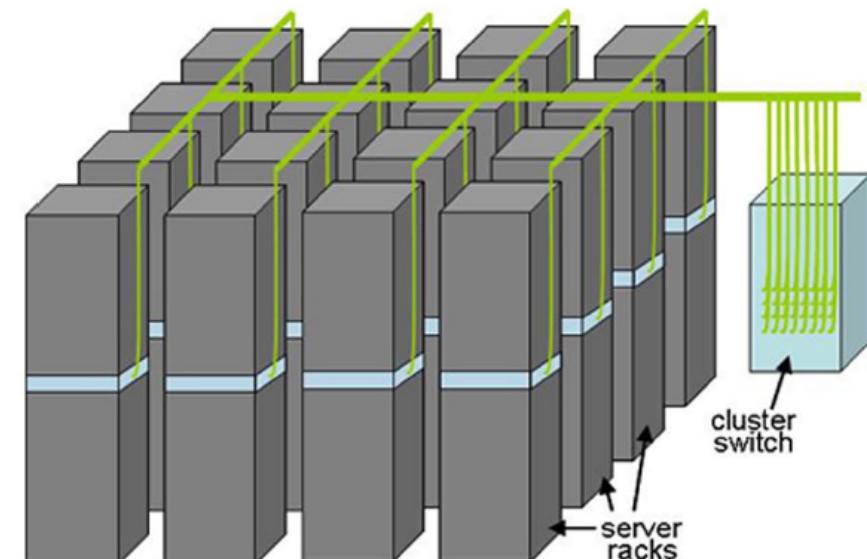
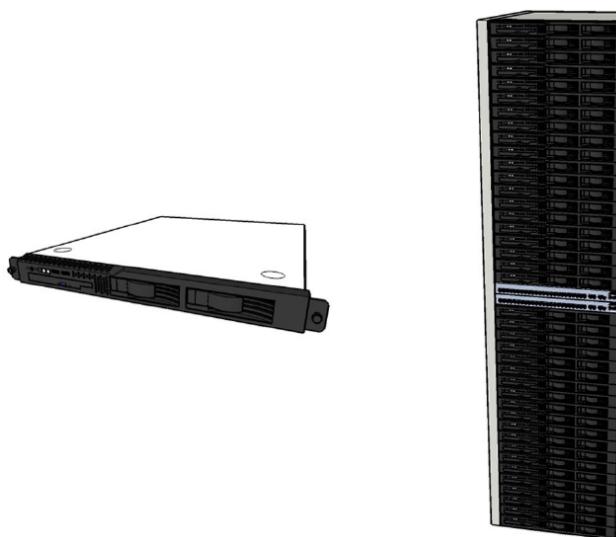
II  
0

# Reliability through replication

---

## 3-way replication

- ▶ Each chunk replicated across 3+ chunkservers
  - ▶ 1 replica in the same rack
  - ▶ 2+ in other, different racks



# Other design decisions

---

Single master to coordinate access

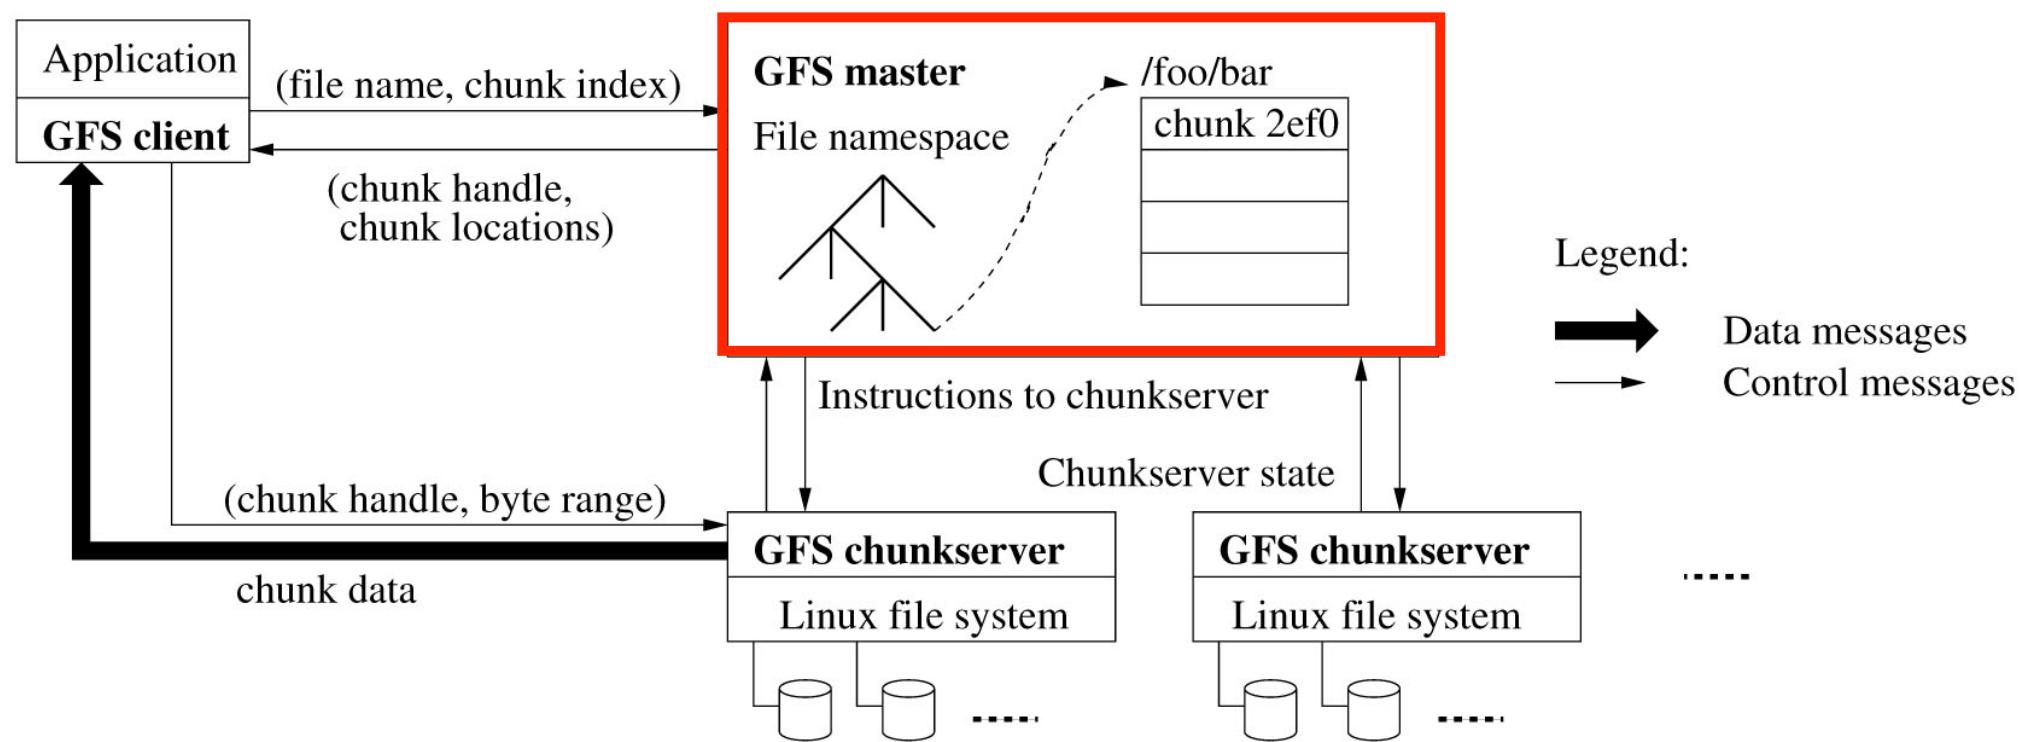
- ▶ keeps *metadata*
  - ▶ *filename, permissions, chunk index, folder hierarchy, etc.*
- ▶ file data stored in chunkservers

Add record append operations

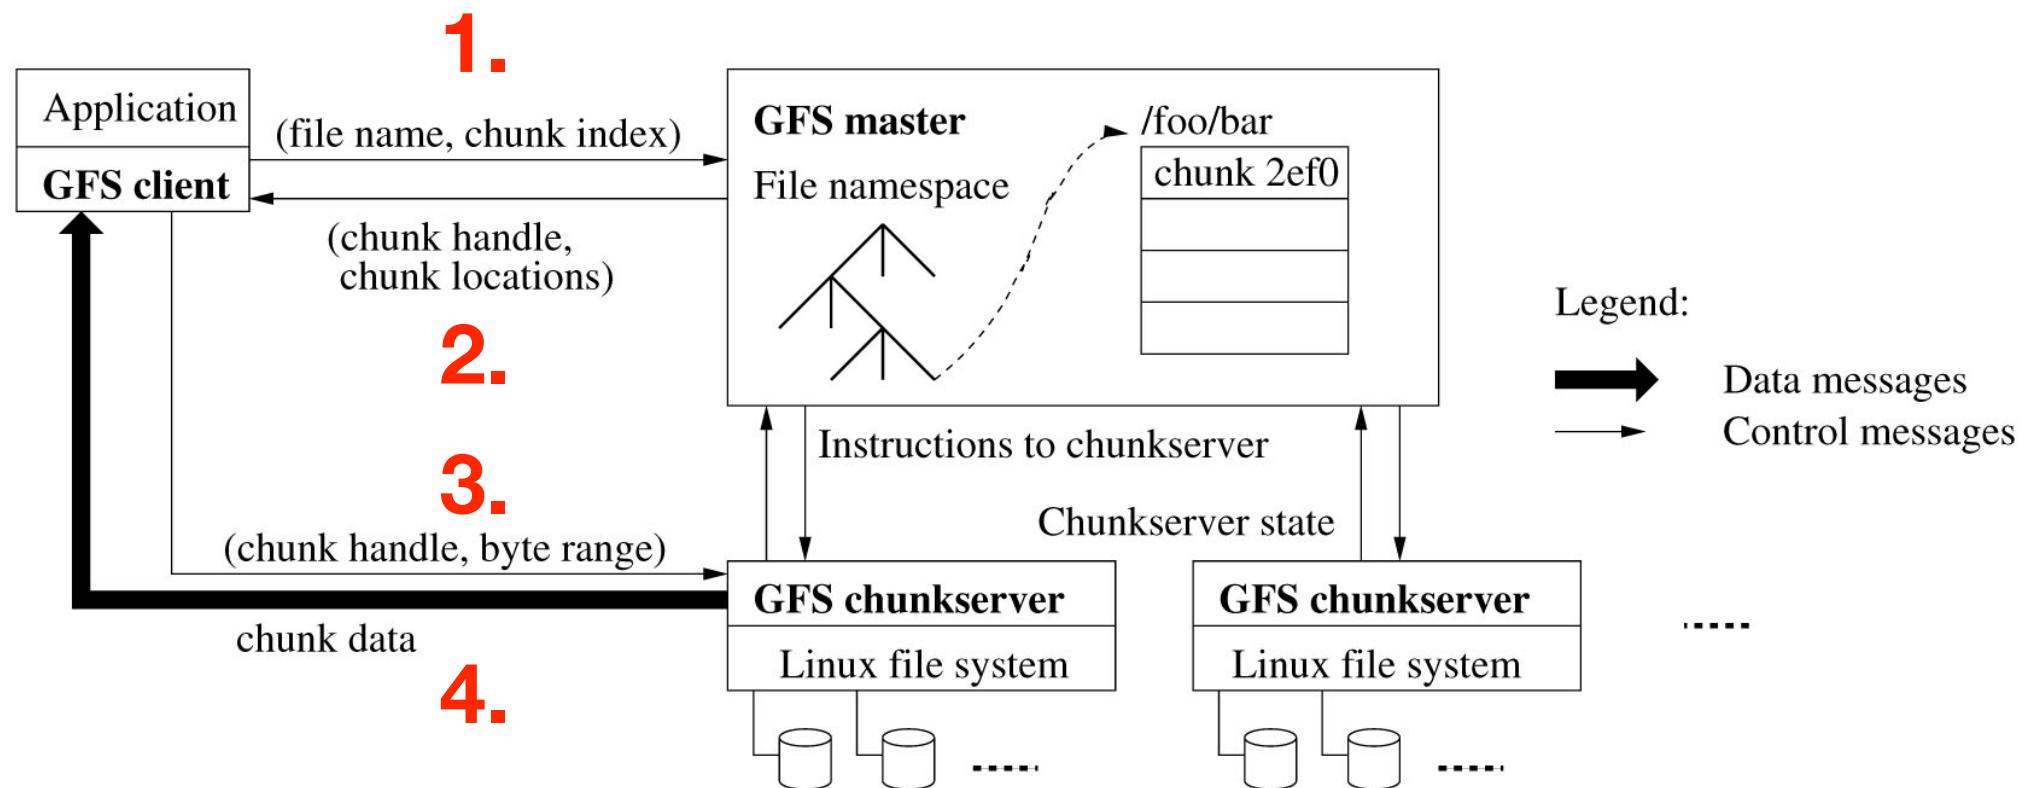
- ▶ Support concurrent appends

# General Architecture

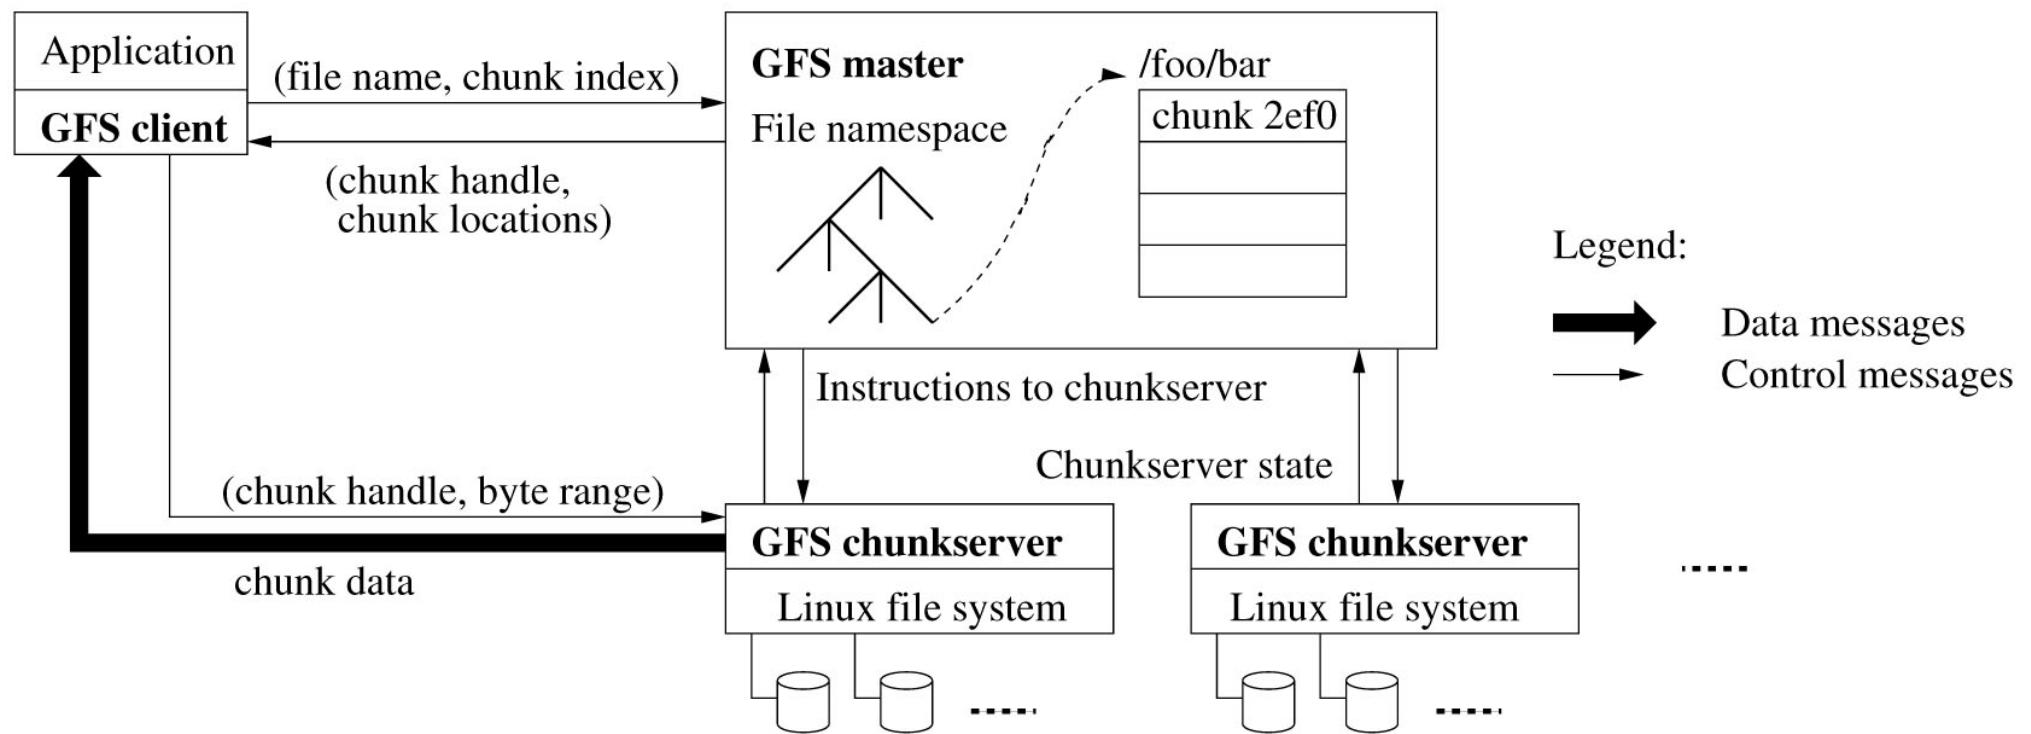
# Single master Multiple chunkservers



# Single master Multiple chunkservers



# Anyone can see the potential problem of this design?



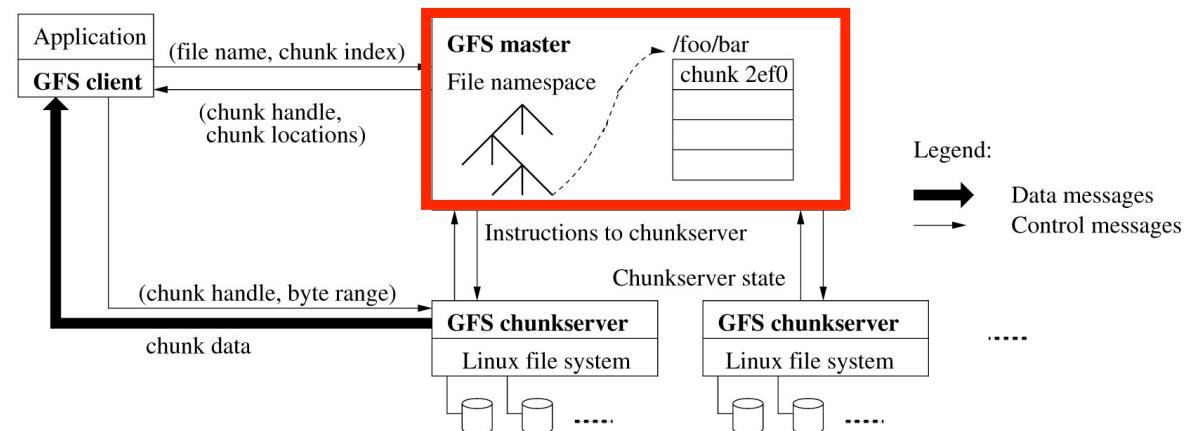
# Single master

---

## Problems and concerns

- ▶ **Single point of failure:** what if master goes offline?
- ▶ **Scalability bottleneck:** what if master is overloaded?

**Any solutions?**



# GFS's answers

---

Master is the single point of failure

- ▶ add a **shadow master**

Master can be overloaded

- ▶ Minimize master involvement to address the scalability issue
  - ▶ Never move data through it, use only for *metadata*
  - ▶ large chunk size: minimizes seeking/indexing time
- ▶ **chunk leases:** master delegates authority to primary replicas in data mutations

# Metadata

---

Three types of metadata, all kept *in memory*

- ▶ file and chunk namespaces
- ▶ mappings from files to chunks (each chunk has a unique ID)
- ▶ locations of each chunk's replicas

# Metadata

---

Three types of metadata, all kept *in memory*

- ▶ file and chunk namespaces
- ▶ mappings from files to chunks (each chunk has a unique ID)
- ▶ locations of each chunk's replicas

First two types are made persistent via an **operation log**

# Metadata

---

Three types of metadata, all kept *in memory*

- ▶ file and chunk namespaces
- ▶ mappings from files to chunks (each chunk has a unique ID)
- ▶ locations of each chunk's replicas

Chunk replica locations learned by polling chunkservers at startup

Chunkserver is final arbiter of what chunks it holds

# Master's responsibilities

---

Metadata storage

Namespace management/locking

Periodic communication with chunkservers

- ▶ give instructions, collect state, track cluster health

Chunk creation, re-replication, rebalancing

- ▶ spread replicas across racks to reduce correlated failures
- ▶ re-replicate data if redundancy falls below a threshold

# Master's responsibilities

---

## Garbage collection

- ▶ simpler, more reliable than traditional file delete
- ▶ master logs the deletion, renames the file to a hidden name
- ▶ lazily garbage collects hidden files

## Stale replica deletion

- ▶ detect “stale” replicas using chunk version numbers

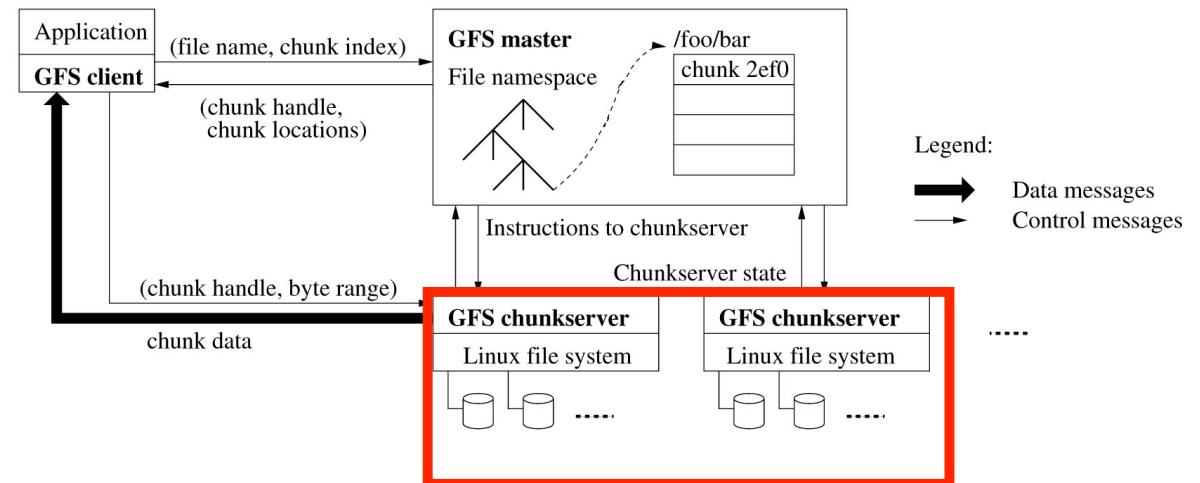
# Chunkserver

Stores 64 MB file chunks on local disk, each with **version number** and **checksum**

Read/write requests specify **chunk handle** and **byte range**

Chunks replicated on configurable number of chunkservers  
(default: 3-way replication)

No file data caching



# Client

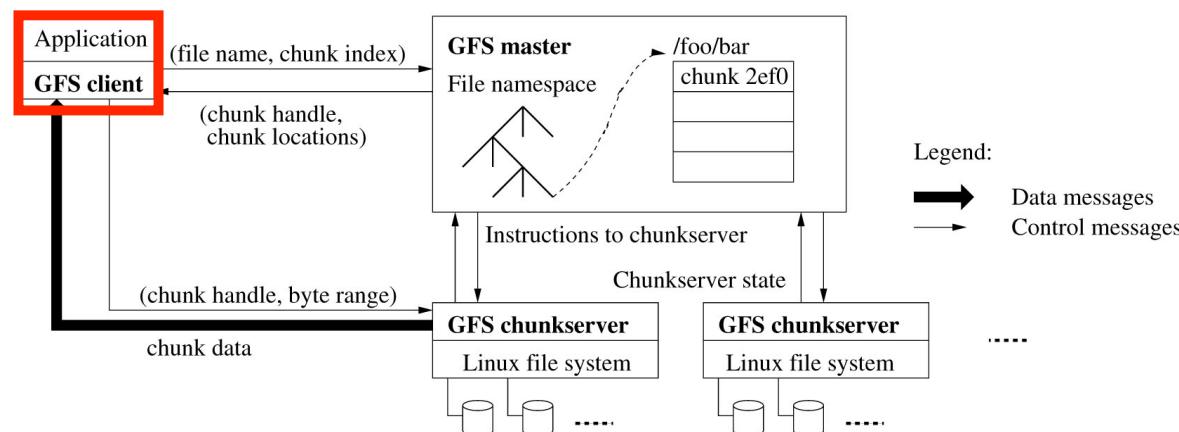
Issues **control (metadata) requests** to master

- ▶ e.g., `ls`

Issues **data requests** directly to chunkservers, e.g., `cat`

- ▶ minimum master involvement

Caches no file data but metadata

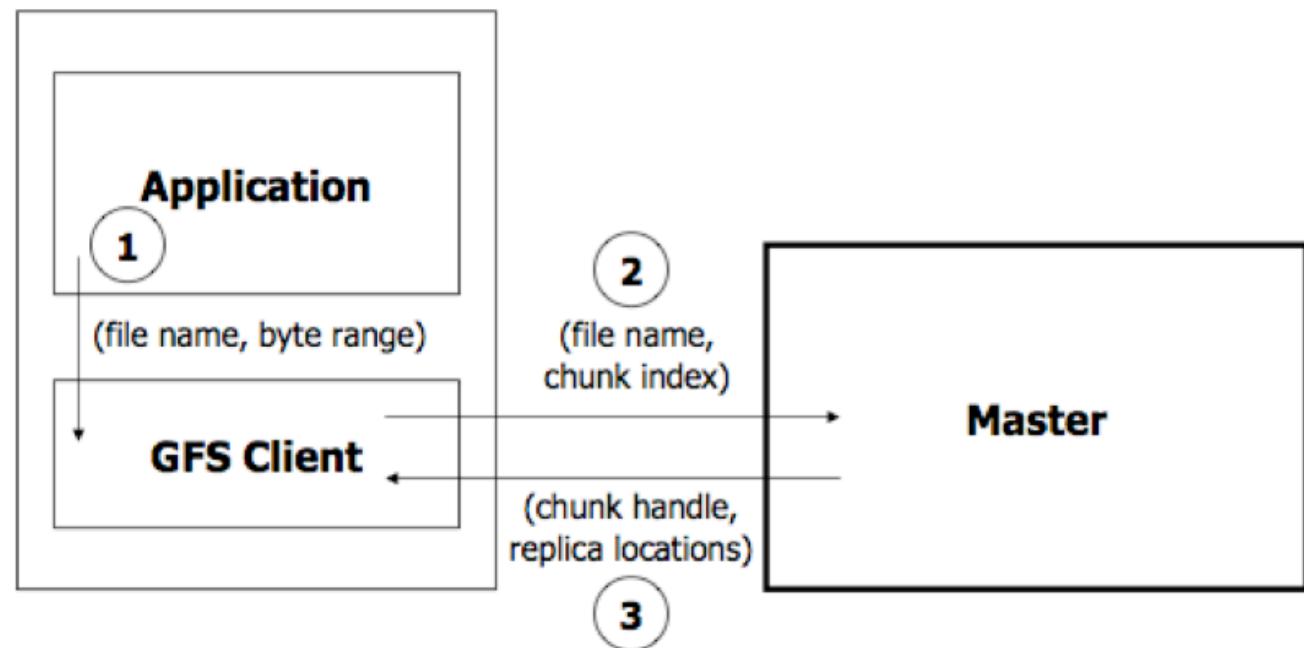


# File read and write

# File read

---

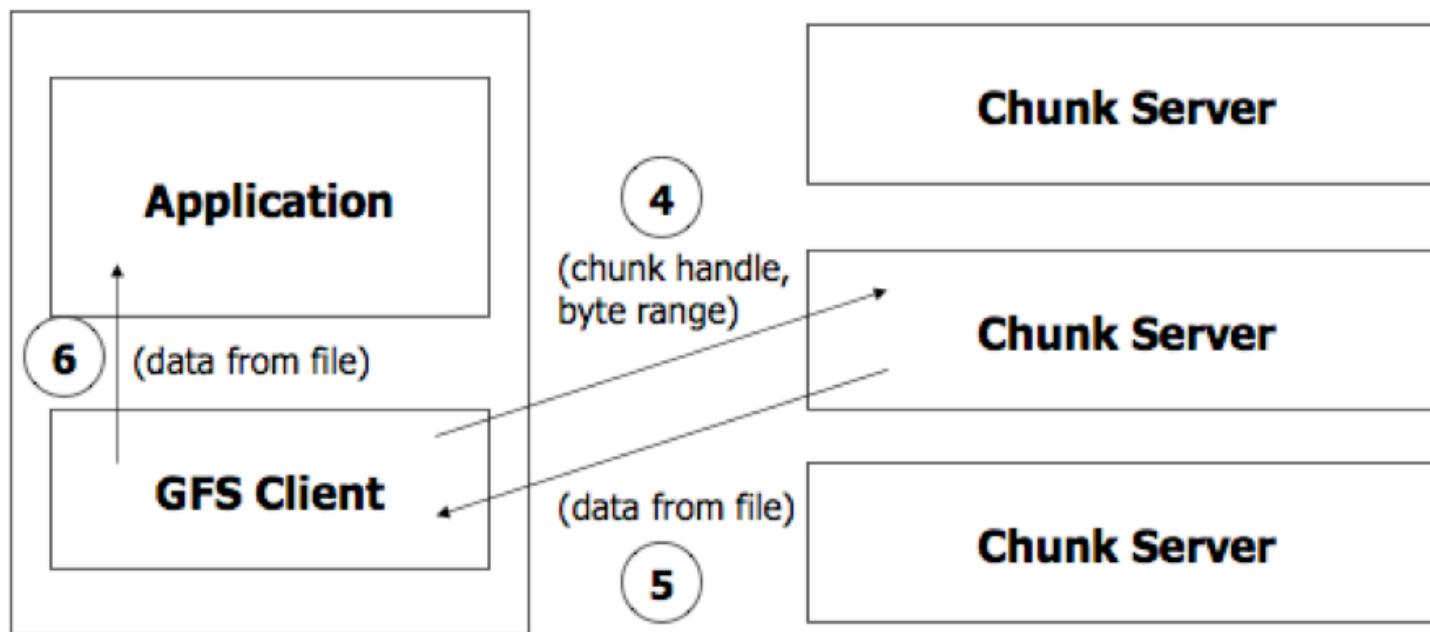
1. Application originates the read request
2. GFS client translates request and sends it to master
3. Master responds with chunk handle and replica locations



# File read

---

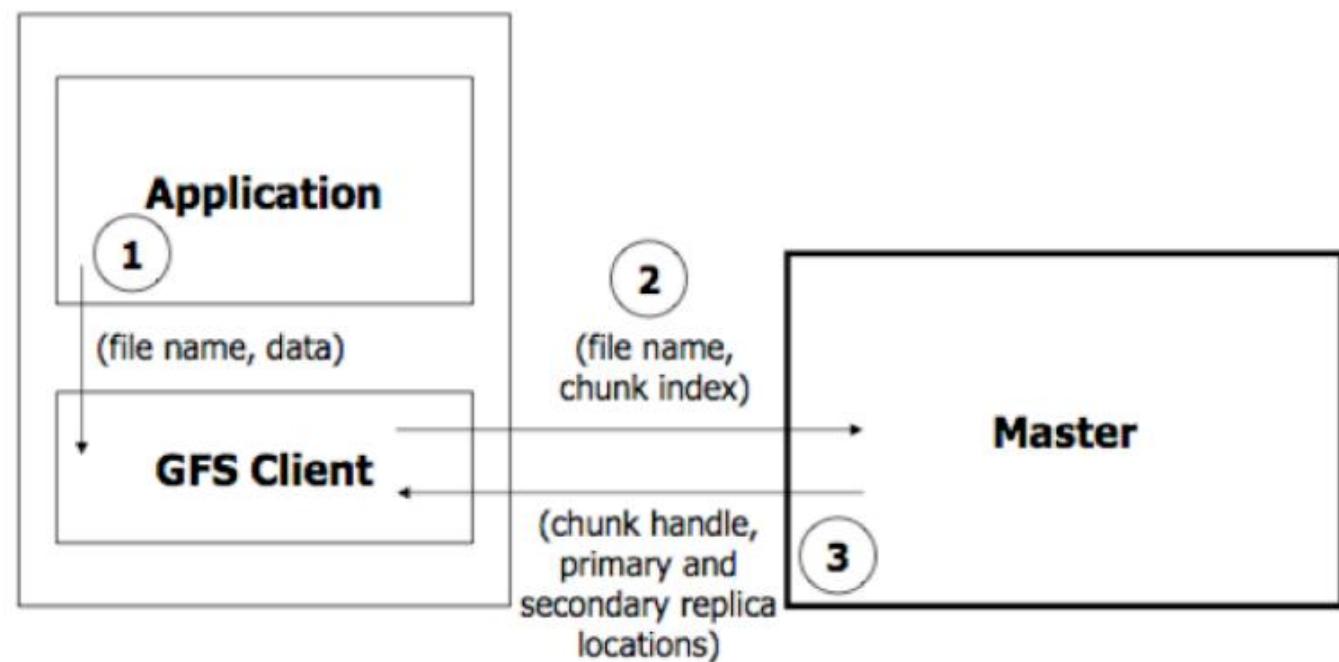
4. Client picks the “closest” location and sends the request
5. Chunkserver sends requested data to the client
6. Client forwards the data to the application



# File write

---

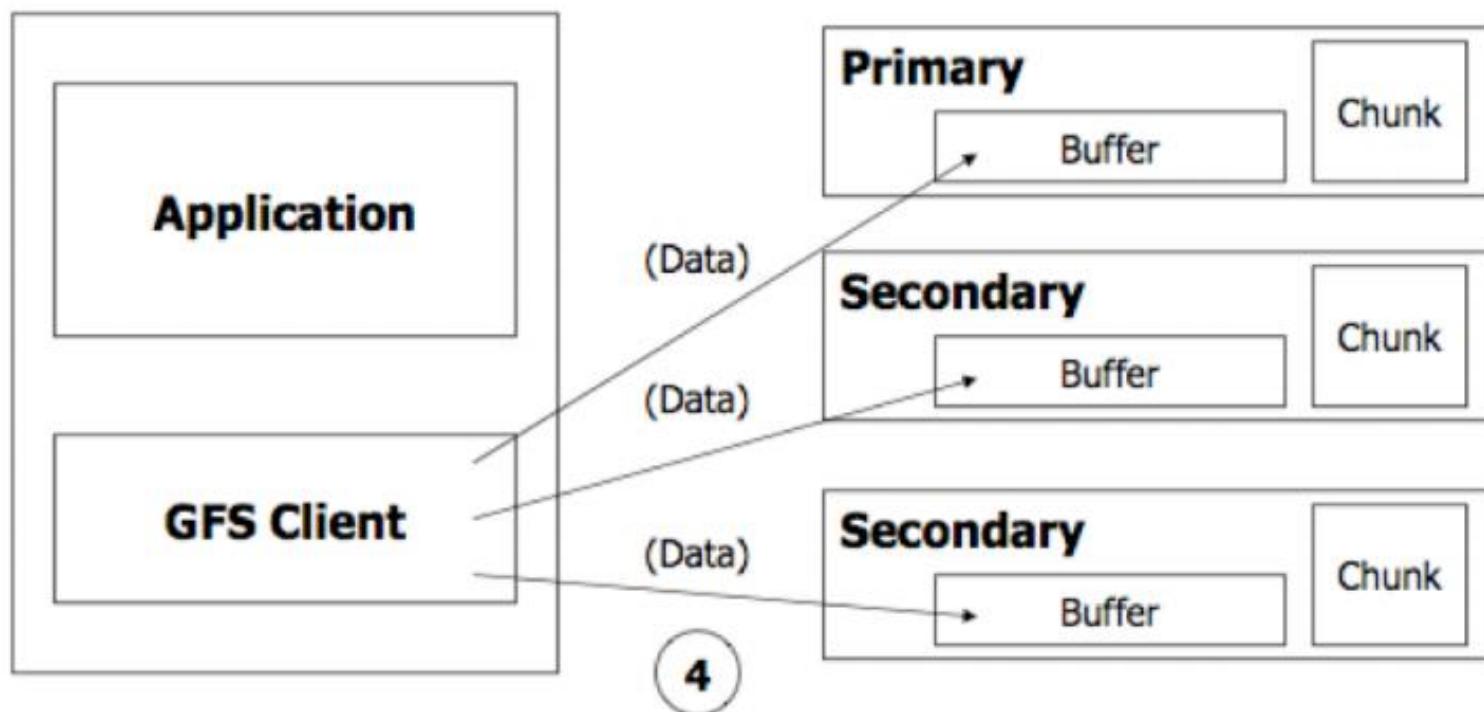
1. Application originates the request
2. GFS client translates request and sends it to master
3. Master responds with chunk handle and replica locations



# File write

---

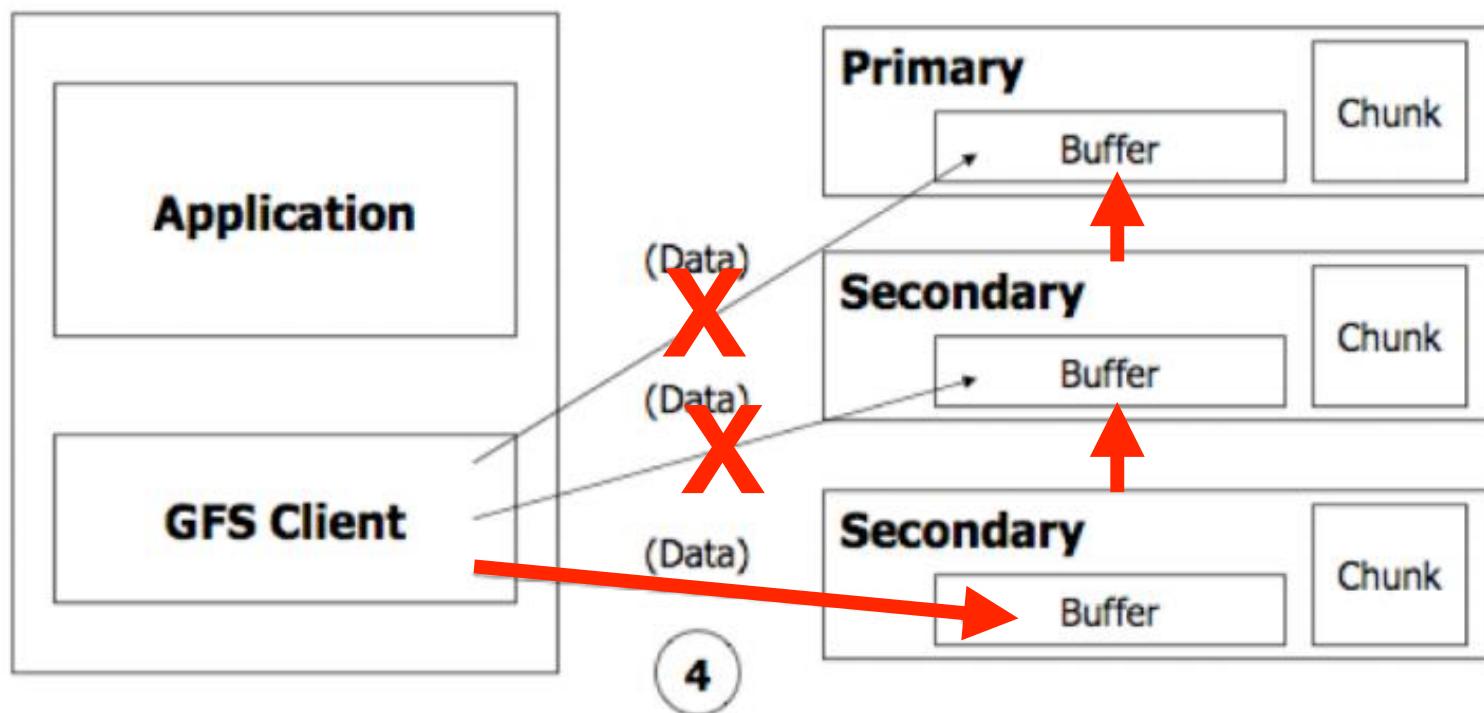
4. Client pushes write data to all locations. Data is stored in chunk server's internal buffers



# File write

4. Client pushes write data to all locations. Data is stored in chunk server's internal buffers

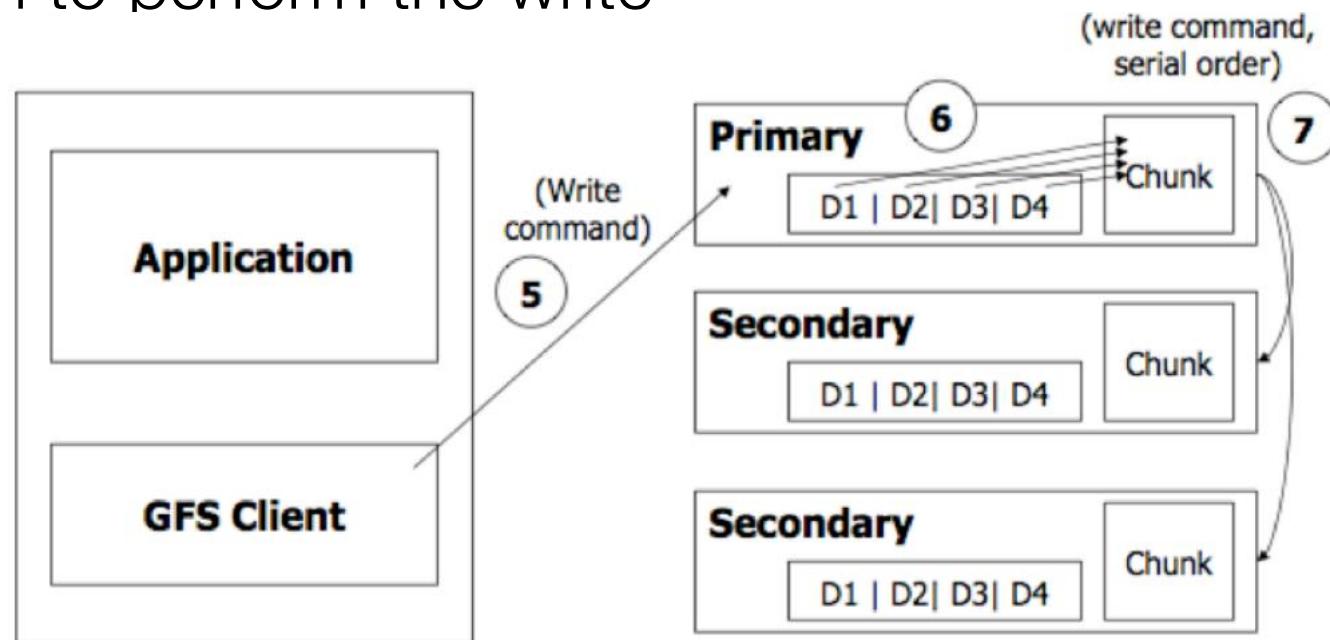
**May form a pipeline**



# File write

---

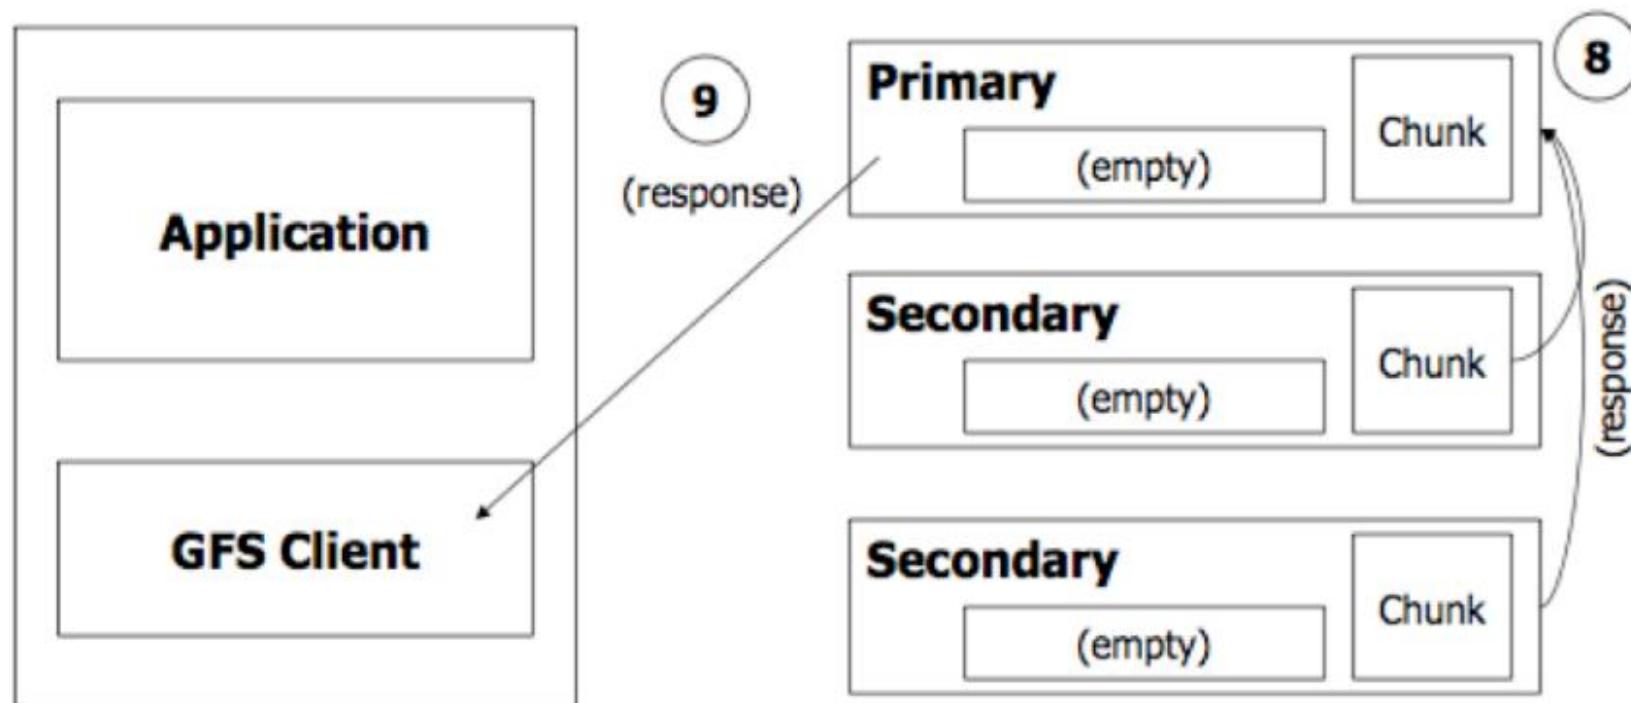
5. Client sends write command to primary
6. **Primary determines serial order for data mutations** in its buffers and writes to the chunk in that order
7. Primary sends the serial order to the secondaries and tells them to perform the write



# File write

---

8. Secondaries respond back to primary
9. Primary responds back to the client



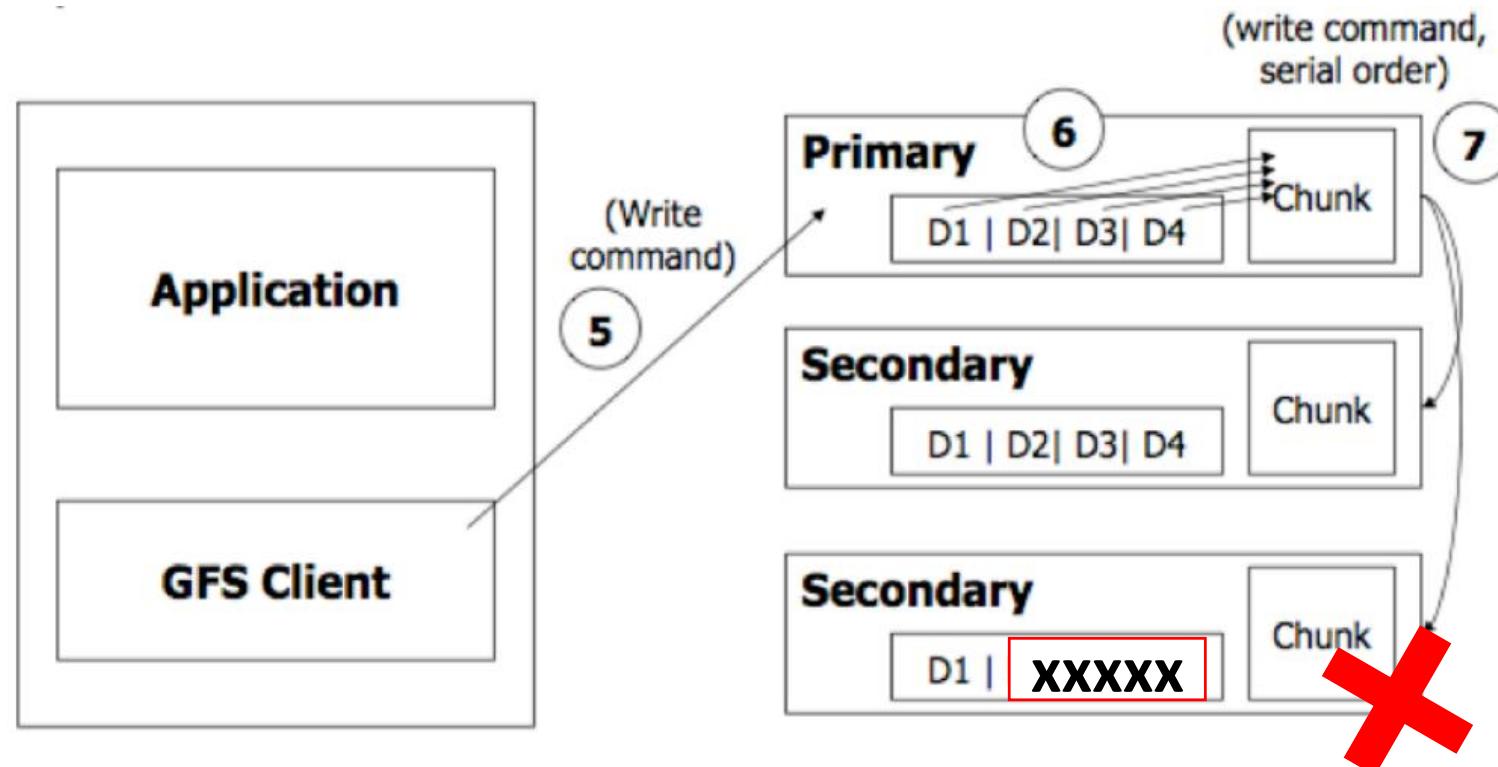
# Fault Tolerance

# What if write fails?

Rewrite

Chunks are not bitwise identical

- ▶ Use checksum to skip inconsistent file regions



# What if chunkserver fails?

---

Master detects a failed “heartbeat” of a chunkserver

Master decrements count of replicas for all chunks on dead chunkserver

Master re-replicates chunks missing replicas elsewhere

# Summary of fault tolerance

---

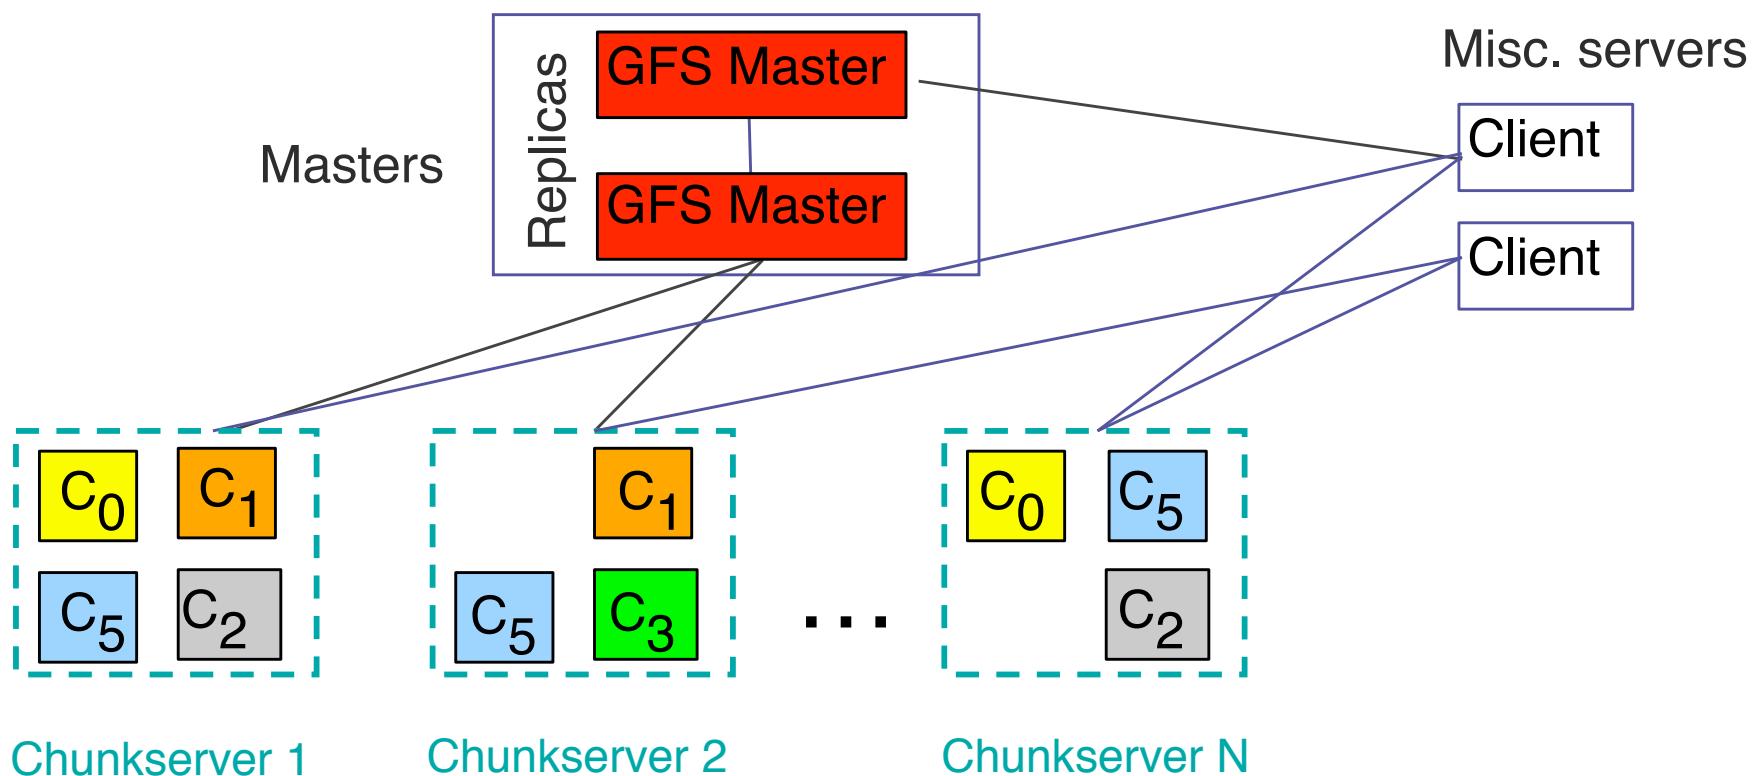
## High availability

- ▶ Fast recovery
  - ▶ master and chunkservers restartable in a few seconds
- ▶ Chunk replication: 3 replicas by default
- ▶ Shadow master

## Data integrity

- ▶ Checksum every 64KB block in each chunk

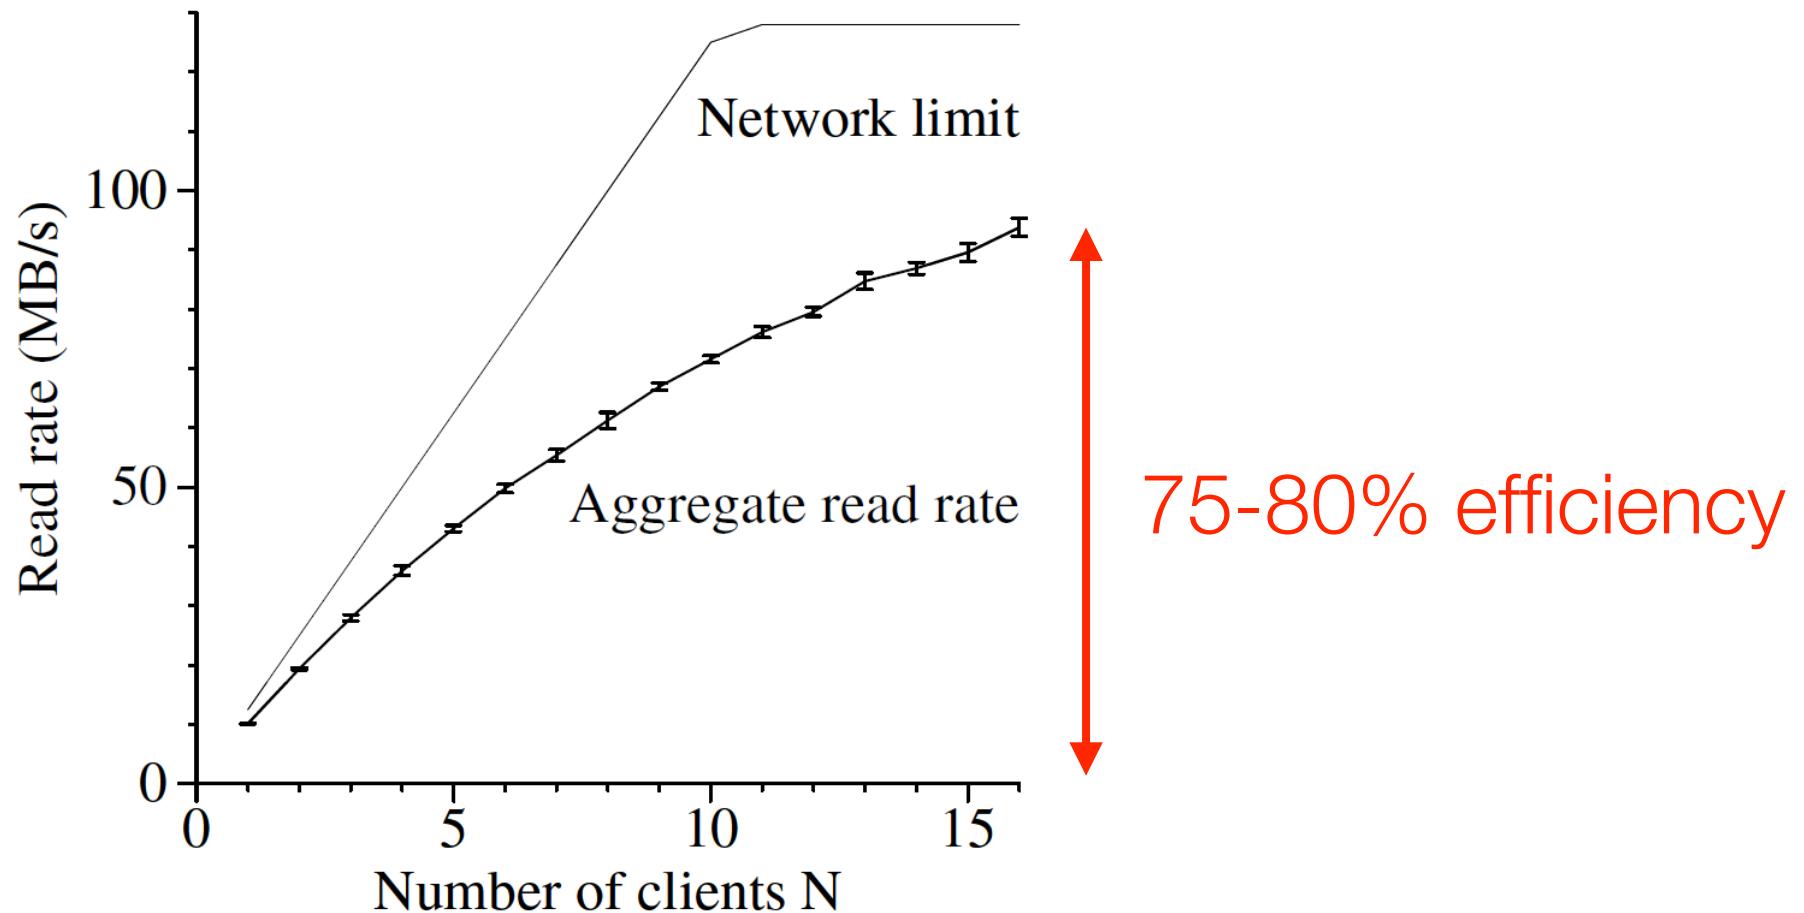
# Summary of fault tolerance



# Measurements (2003)

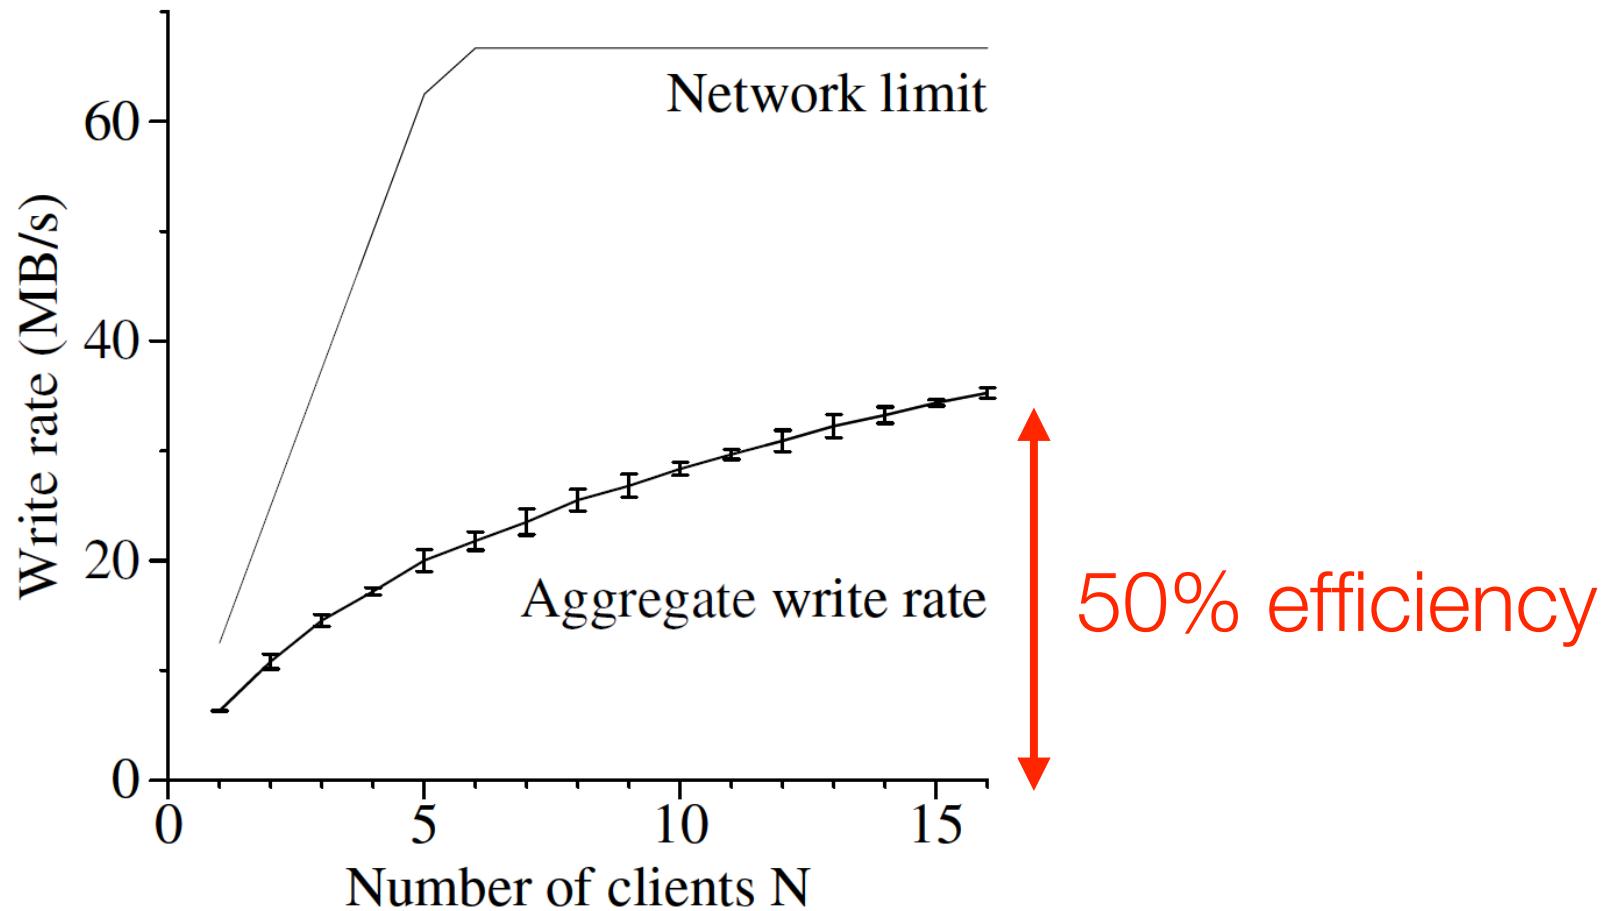
# Micro-benchmark: Reads

---



# Micro-benchmark: Writes

---



# Real cluster: Recovery time

---

## Kill 1 chunkserver

- ▶ 15,000 chunks (= 600 GB of data) restored in 23.2 mins
- ▶ effective replication rate = 440 MB/s

## Kill 2 chunkservers

- ▶ 16,000 chunks on each (= 660 GB of data)
- ▶ reduced 266 chunks to having a single replica
- ▶ 266 chunks restored to at least 2x replications within 2 mins

# Conclusions

---

GFS demonstrates how to support large-scale processing workloads on commodity hardware

- ▶ tolerates frequent component failures (failures as the norm rather than the exception)
- ▶ optimizes for **huge** files that are mostly **appended** and then **read sequentially**
- ▶ delivers high aggregate throughput to many concurrent readers and writers

# Any limitations?

# Limitations

---

Assumes write-once, read-many workloads

Assumes a modest number of large files

- ▶ **Large chunk size:** small files can create hot spots on chunkservers if many clients accessing the same file

# Credit

---

Some slides are adapted from

- ▶ Dr. Lazowska and Dr. Chernyak's slides for CSE 490H at UW
- ▶ Dr. Alex Moshchuk's slides used in the Google lecture series at UW
- ▶ Dr. Jeff Dean's slides used in the Google lecture series at UW
- ▶ Dr. Romain Jacotin's slides